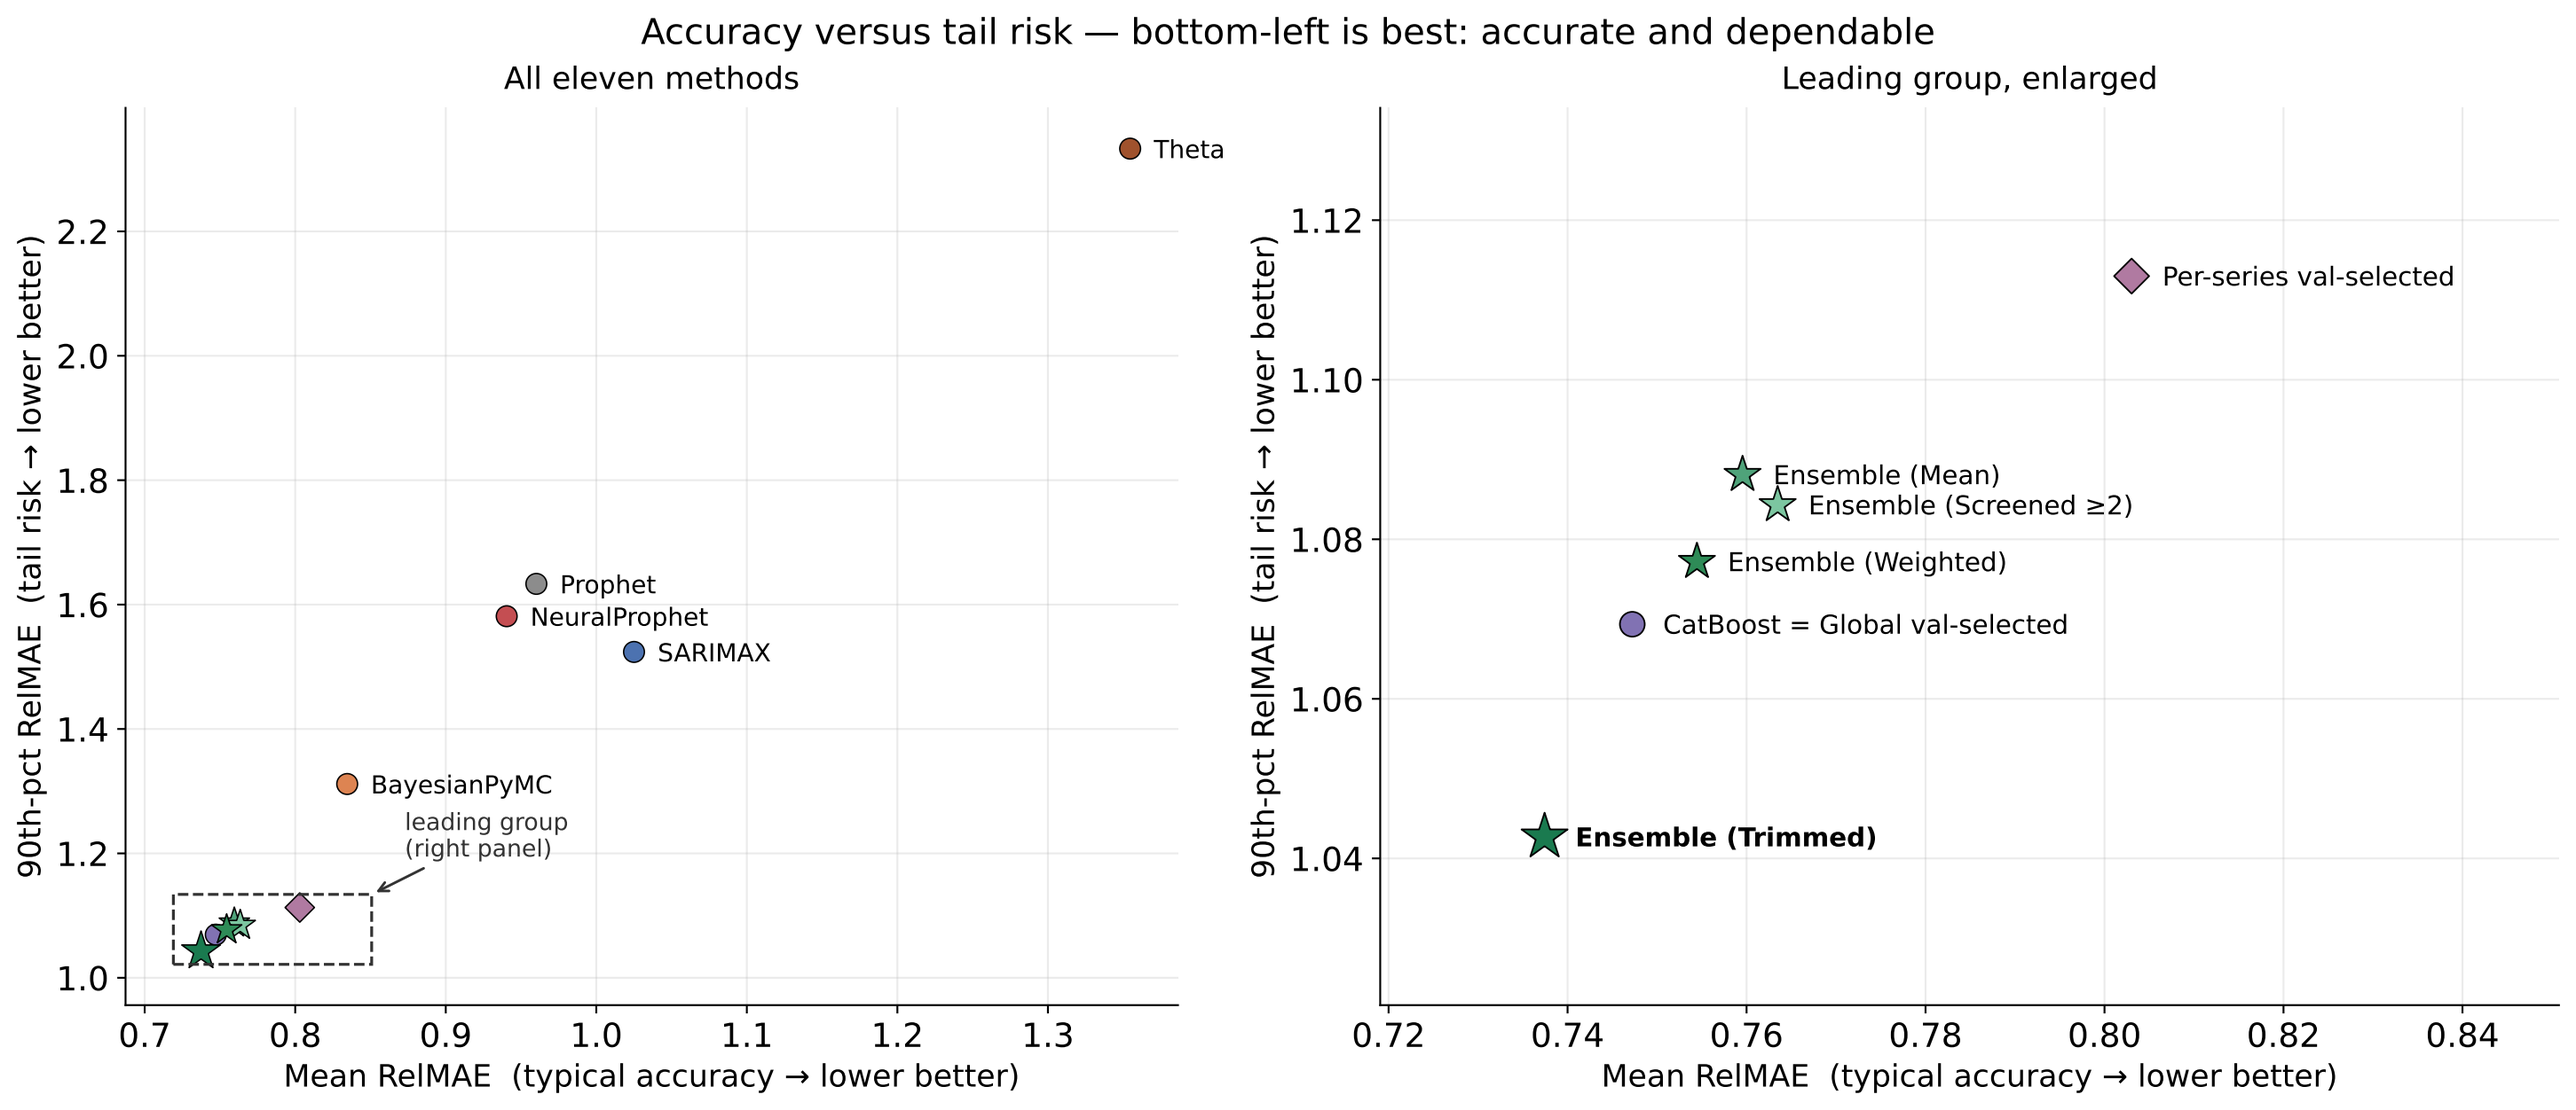
Accuracy versus tail risk — bottom-left is best: accurate and dependable
All eleven methods
Leading group, enlarged
Theta
Prophet
NeuralProphet
SARIMAX
BayesianPyMC
leading group
(right panel)
Per-series val-selected
Ensemble (Mean)
Ensemble (Screened ≥2)
Ensemble (Weighted)
CatBoost = Global val-selected
Ensemble (Trimmed)
90th-pct RelMAE (tail risk → lower better)
Mean RelMAE (typical accuracy → lower better)
1.0
1.2
1.4
1.6
1.8
2.0
2.2
0.7
0.8
0.9
1.0
1.1
1.2
1.3
1.04
1.06
1.08
1.10
1.12
0.72
0.74
0.76
0.78
0.80
0.82
0.84